

Customer 360 Next Best Action

Turn scattered service signals into a clear, evidence-led next step.

A working local prototype for insurance and lending service teams.

DEMO DATA IN SNOWFLAKE

10

customers

12

accounts

11

interactions

10

Customer 360 rows

Problem brief: fragmented records slow service recovery

THE OPERATIONAL GAP

Agents piece together customer context across products, cases, and call notes.

01 Context is split

Profile, account, and interaction history sit in different views.

02 Urgency is easy to miss

An unresolved claim or billing issue can be buried in a long history.

03 Actions vary by agent

Teams lack a shared reason, owner, and next-step playbook.

WHO NEEDS A BETTER HANDOFF

● Claims operations lead

● Support specialist


● Relationship manager

INDUSTRY CONTEXT


Insurance · Lending · Banking
Claims, billing, renewals, and service operations

Architecture: one verified warehouse view, one local app

VERIFIED SNOWFLAKE DATA PATH



CURRENT APP PATH · SEPARATE FROM SNOWFLAKE



Future: Cortex, governance, and app-to-view connection. CoCo verified: sql-author query + data-quality audit (no duplicate IDs, null keys, or orphans).

Evidence-led service recovery, with a human in control

01

Signal

Open claim delay

02

Evidence

I5001 transcript

03

Decision

Fast-track review

04

Follow-through

Agent records outcome

CASE C1001 · OPEN CLAIM DELAY

"The hospital claim has been pending for three weeks... support promised a callback twice."

Source: interaction I5001 · supplied label: Negative · status: Open

RULE OUTPUT · NOT AN AI CLAIM

Risk score 65 / 100

P1 · same-day action
Claims operations lead

The score is a transparent service-risk heuristic, not a calibrated churn probability. No message is sent automatically.

Impact: a measurable hypothesis, not a claimed result

WHAT THE DEMO CAN SHOW TODAY

10

synthetic customers in the unified view

5

open interaction records in the demo queue

10

automated regression tests passing

NEXT VALIDATION STUDY

1

Compare

Manual record review vs. evidence-led triage

2

Measure

Median triage time and evidence correctness

3

Pilot

Case aging, SLA resolution, and agent adoption

Scale path: incremental Snowflake ingestion → governed semantic layer → bounded Cortex enrichment → connected Streamlit app → outcome evaluation.